

A Generic Communication Scheduler for Distributed DNN Training Acceleration

Yanghua Peng, Yibo Zhu, Yangrui Chen, Yixin Bao, Bairen Yi, Chang Lan,
Chuan Wu, Chuanxiong Guo

The University of Hong Kong



ByteDance Inc.



DNN Training



DNN training is compute-hungry and time-consuming

ResNet50	Training Time	BERT	Training Time
1 TPUv3	10 hours	16 TPUv3	81 hours
1024 TPUv3	1.28 minutes	1024 TPUv3	76.19 minutes


<https://mlperf.org/training-results-0-6>

<https://arxiv.org/pdf/1904.00962.pdf>


Training can be scaled out by data parallelism or model parallelism



Data Parallel DNN Training



PS Dependency Graph




Dependency:

- Backward depends on forward
- Push depends on backward
- Pull depends on push
- Forward depends on pull

Framework engines execute the graph according to the dependencies

Communication Scheduling




40% improvement

Problem: FIFO strategy does not overlap communication with computation well

P3, TicTac: partition tensors and change tensor transmission order

Communication Scheduling



Problem: FIFO strategy does not overlap communication with computation well

P3, TicTac: partition tensors and change tensor transmission order

Limitations of Existing Work


P3 and TicTac:

- Coupled with **specific** framework implementations, e.g., P3 for MXNet PS and TicTac for TensorFlow PS.
- Heuristic scheduling with empirical results

Many different setups in distributed DNN training:



P Y T Ø R C H



Communication architectures



Network protocols

ML frameworks

Limitations of Existing Work

P3 and TicTac:


- Coupled with **specific** framework implementations, e.g., P3 for MXNet PS and TicTac for TensorFlow PS.
- ~~Heuristic scheduling with empirical results~~

Many different setups in distributed learning

How to do communication scheduling:

1. Work in all setups
2. Minimal modifications
3. Scheduling Optimality

P Y T Ø R C H



R D M A

ML frameworks

Communication architectures

Network protocols

One Unified Scheduling System for ALL


Observation: The dependency graph structure is [intrinsic](#) for DNN training, regardless of training frameworks, communication architectures, or network protocols

ByteScheduler: A [generic](#) tensor scheduling framework

- One unified scheduler framework that abstracts tensor scheduling from various frameworks, communication architectures and network protocols
- One principled scheduling algorithm that is guided by theory and works in real-world

Challenge 1: Different ML Frameworks



- Imperative framework (e.g., PyTorch) and declarative framework (e.g., TensorFlow)
- Global barrier between iterations (e.g., TensorFlow, PyTorch), causing any scheduling of push/all-reduce ineffective



Challenge 2: Different Runtime Environments

The overhead of scheduling & tensor partitioning is different for different system setups and network conditions

How to balance the performance gain with scheduling overhead? The system parameters (e.g., partition size) are likely to be affected by different runtime configurations, e.g., bandwidths, DNNs




Outline

1. Background and Motivation

2. ByteScheduler Design

3. Evaluation

Unified Scheduler Across Frameworks




Unified Scheduler Across Frameworks

DNN definition

ByteScheduler

Graph execution


Communication



ByteScheduler Architecture

Plugin: Wrap each communication operation as a CommTask

Core: Partition and schedule CommTasks



CommTask: A Unified Abstraction

CommTask: A wrapped communication operation, e.g., push one tensor, all-reduce one tensor

CommTask APIs implemented in framework plugins:

- `partition(size)`: partition a CommTask into SubCommTasks with tensors no larger than a threshold `size`
- `notify_ready()`: notify Core about the readiness of a CommTask
- `start()`: start a CommTask by calling the underlying push/pull/all-reduce
- `notify_finish()`: notify Core about the completion of a CommTask


Dependency Proxy: Get the Scheduling Control

A **Dependency Proxy** is an operator to get the scheduling control from the frameworks to the Core

Dependency Proxy:

- Trigger CommTask.notify_ready() via a callback
- Wait to finish until Core calls CommTask.start()
- Generate completion signal using CommTask.notify_finish()


Implementation differs for
imperative and declarative engines



Dependency Proxy: Crossing the Global Barrier

Out-of-engine communication: Start the actual communication outside engine


Layer-wise out-of-engine dependencies: Build correct dependency for each layer by adding a Proxy to block forward computation



Dependency Proxy: Crossing the Global Barrier

Out-of-engine communication: Start the actual communication outside engine

Layer-wise out-of-engine dependencies: Build correct dependency for each layer by adding a Proxy to block forward computation



Optimal Scheduling Theorem

Optimal scheduling for minimizing the time for each iteration:

- For PS, prioritize $push_i$ over $push_j$, and $pull_i$ over $pull_j$, $\forall i < j$
- For all-reduce, prioritize $allreduce_i$ over $allreduce_j$, $\forall i < j$
- Assuming infinitely small partition size and immediate preemption without overhead

In practice, partitioning and preemption have overhead

Credit-based Preemption

Stop-and-wait approach in previous work can not fully utilize network bandwidth

- Send a single tensor and wait for its ACK

Credit-based Preemption

- Work like a sliding window and the credit is the window size
- Allow **multiple** tensors in a sliding window to be sent concurrently

Credit size is an important system parameter


- Pro: higher bandwidth utilization
- Con: less timely preemption due to FIFO communication stack

Auto-tuning Partition Size and Credit Size

Optimal partition size and credit size are affected by many factors, e.g., network bandwidths, number of workers, DNN models, CPU and GPU types

We use [Bayesian Optimization](#) for auto-tuning

- Work with general objective function
- Minimize the overhead, i.e., the number of sampled points



Outline

1. Background and Motivation

2. ByteScheduler Design

3. Evaluation

Evaluation

Implementation: MXNet PS and all-reduce (based on Horovod), PyTorch (based on Horovod), TensorFlow PS


```
# After user created an MXNet KVStore object kvs
from bytescheduler.mxnet.kvstore import ScheduledKVS
kvs = ScheduledKVS(kvs)
# Continue using kvs without any further modification
```

Testbed: 16 machines, each with 8 Tesla V100 GPUs and 100Gbps bandwidth


Comparison:

- Baseline: vanilla ML frameworks
- Linear scaling: vanilla training speed on 1 machine multiplied by the number of machines


Scalability of ByteScheduler



VGG16 (97%-125%)



ResNet50 (9%-15%)




Transformer (70%-171%)


MXNet PS RDMA

- Up to 171% improvement and close to linear scaling


ByteScheduler for Multiple Frameworks



MXNet PS RDMA
(97%-125%)



TensorFlow PS RDMA
(170%-196%)




PyTorch NCCL TCP
(7%-13%)


VGG16

- Up to 196% improvement compared to the baseline


ByteScheduler Adapts to Different Bandwidths



VGG16 (79%-132%)



ResNet50 (10%-64%)



Transformer (67%-70%)

MXNet PS RDMA

- Consistent speedup in all bandwidth settings
- Without auto-tuning, the training speed is lower

Conclusion

ByteScheduler: A generic communication scheduler for distributed DNN training acceleration

- Unified abstraction for tensor scheduling
- Multiple training framework support, with minimal code change to existing frameworks
- Principled tensor scheduling design with parameter autotuning

Q&A

Source code:

<https://github.com/bytedance/byteps/tree/bytescheduler/bytescheduler>