


# Verification of Software Network Functions with No Verification Expertise

Arseniy Zaostrovnykh, Solal Pirelli, Rishabh Iyer, Matteo Rizzo,  
Luis Pedrosa, Katerina Argyraki, George Canea



# Bright Future With Verification



# Bright Future With Verification



Push-button

Pay-as-you-go

Full-stack

Vigor enables practical verification of software network functions.

- Virtualized networks
- Critical infrastructure, but unreliable software
- NF operators and devs have no verification expertise

Push-button

Pay-as-you-go

Full-stack

No human effort, nor verification expertise

- Traditional verification cost: 5-10x development effort<sup>†</sup>
- Ultimate goal: zero-cost verification

<sup>†</sup>seL4 [SOSP 2009]  
Ironclad [OSDI 2014]  
HACL\* [CCS 2017]

Push-button

Pay-as-you-go

Full-stack

Support for partial specifications

EXT\_PORT = 2

```
if received_on_port != EXT_PORT: # Heartbeat
    return ([], [])
```

```
1  from state import flow_emap, flow_id_to_backend_id, backends, backend_ip_emap, cht
2  EXP_TIME = 10 * 1000
3  BACKEND_EXP_TIME = 360000000 * 1000
4  EXT_PORT = 2
5
6  if a_packet_received:
7      flow_emap.expire_all(now - EXP_TIME)
8      backend_ip_emap.expire_all(now - BACKEND_EXP_TIME)
9
10 h3 = pop_header(tcpudp, on_mismatch=([], []))
11 h2 = pop_header(ipv4, on_mismatch=([], []))
12 h1 = pop_header(ether, on_mismatch=([], []))
13
14 assert a_packet_received
15 assert h1.type == 8 # 0x0800 == IPv4 in big endian
16 assert h2.npid == 6 or h2.npid == 17 # 6/17 -> TCP/UDP
17
18 if received_on_port == EXT_PORT: # Packet from the external network - client
19     packet_flow = LoadBalancedFlow(h2.saddr, h2.daddr, h3.src_port, h3.dst_port, h2.npid)
20     alloc_flow_and_process_packet = False;
21     if flow_emap.has(packet_flow):
22         flow_id = flow_emap.get(packet_flow)
23         backend_id = flow_id_to_backend_id.get(flow_id)
24         if backend_ip_emap.has_idx(backend_id):
25             flow_emap.refresh_idx(flow_emap.get(packet_flow), now)
26             backend = backends.get(backend_id)
27             return ([backend.nic],
28                     [ether(h1, saddr=..., daddr=backend.mac),
29                      ipv4(h2, cksum=..., daddr=backend.ip),
30                      tcpudp(h3)])
31     else:
32         flow_emap.erase(packet_flow)
33         alloc_flow_and_process_packet = True
34
35     else:
36         alloc_flow_and_process_packet = True
37         if alloc_flow_and_process_packet:
38             if backend_ip_emap.exists_with_cht(cht, _LoadBalancedFlow_hash(packet_flow)):
39                 bknd = backend_ip_emap.choose_with_cht(cht, _LoadBalancedFlow_hash(packet_flow))
40                 if not flow_emap.full():
41                     idx = the_index_allocated
42                     flow_emap.add(packet_flow, idx, now)
43                     flow_id_to_backend_id.set(idx, bknd)
44                     backend = backends.get(bknd)
45                     return ([backend.nic],
46                             [ether(h1, saddr=..., daddr=backend.mac),
47                              ipv4(h2, cksum=..., daddr=backend.ip),
48                              tcpudp(h3)])
49             else:
50                 return ([], [])
51         else: # A heartbeat from a backend
52             bknd_addr = ip_addr(h2.saddr)
53             if backend_ip_emap.has(bknd_addr):
54                 backend_ip_emap.refresh_idx(backend_ip_emap.get(bknd_addr), now)
55             else:
56                 if not backend_ip_emap.full():
57                     idx = the_index_allocated
58                     backend_ip_emap.add(bknd_addr, idx, now)
59                     backends.set(idx, LoadBalancedBackend(received_on_port, h1.saddr, h2.saddr))
60
61     return ([], [])
```

Push-button

Pay-as-you-go

Full-stack

Whole SW stack: framework, OS, drivers

NF: 3 **KLOC**

Framework, driver: 85 **KLOC**



Push-button

Pay-as-you-go


Full-stack

Whole SW stack: framework, OS, drivers

NF: 3 **KLOC**

Framework, driver: 85 **KLOC**

Linux: 23 **MLOC**



**Vigor enables**

Push-button ,

Pay-as-you-go , and

Full-stack **verification of software NFs.**

# Outline

Push-button

Pay-as-you-go

Usage scenario


Full-stack

Evaluation [ Practical | No performance overhead ]


[vigor.epfl.ch](http://vigor.epfl.ch)

**Network Function (NF)** is  
the core of a network middlebox


# NF in a Nutshell



# NF = NF Logic + Data Structs



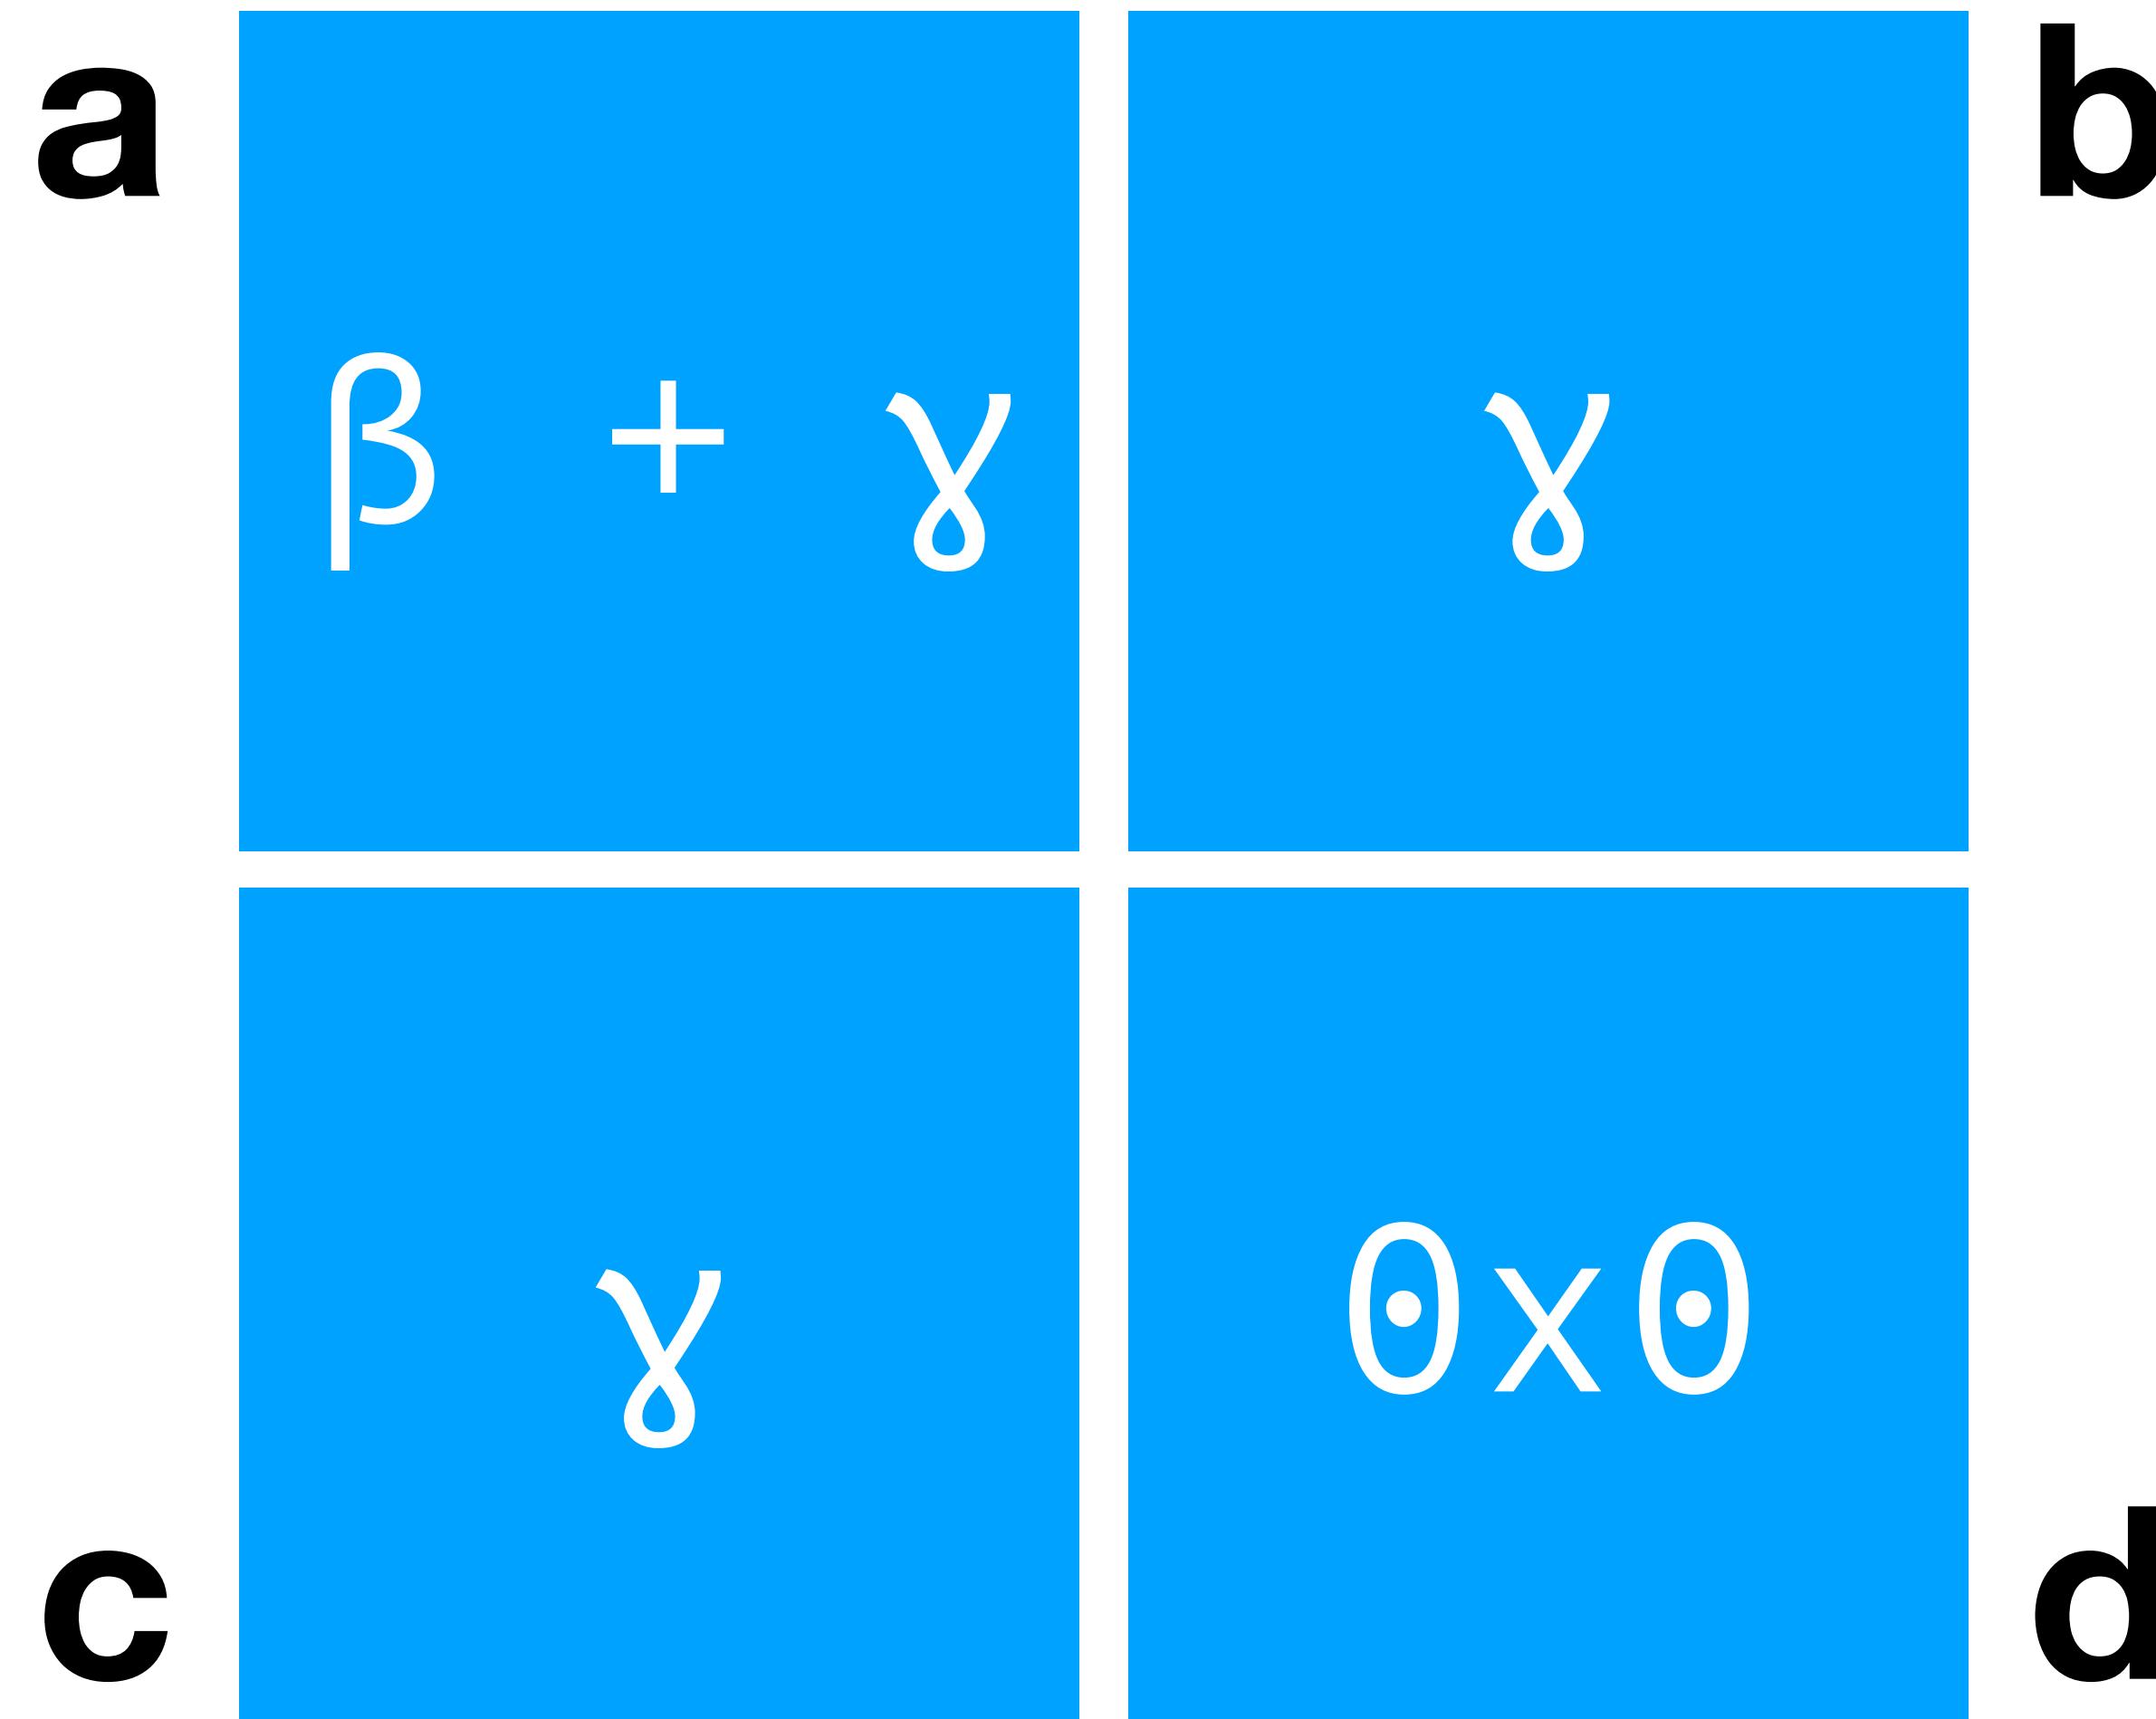
# 2 Verification Approaches



# **SE = Low-Level & Automatic**

**a**  $\leftarrow$  **b** + **c**


$\Rightarrow$  **b**  $\leftarrow$  **c** + **d**



# TP = High-Level & Manual

**a**  $\leftarrow$  \***b**

$\Rightarrow$  **b**  $\leftarrow$  \*(**b** + 1)



# TP = High-Level & Manual


*/\*@ a - head, b - tail @\*/*

**a** <- \***b**


*/\*@ This is replacing  
the list head @\*/*

➤ **b** <- \*(**b** + 1)


*/\*@ This is replacing  
the list tail @\*/*




# Verification Steps




# Semantic Check



# Semantic Check



# Who Writes The Specs



# Specification Abstraction

**Where there's a will, there's a way.**

Vouloir, c'est pouvoir.

Wo ein Wille ist, ist auch ein Weg.

Было бы желание, а способ найдётся.

# Common Abstractions

 Party	Abstractions used
IETF (NAT, Firewall)	same data structs
IEEE (bridge)	same data structs
In-house design (Maglev)	same data structs
Developers	same data structs

Common abstractions do exist: NF data structures


# Insight

**NF implementations and  
specifications use the  
same set of abstractions**

# Verification Effort

5-10x

# Verification Effort



**0.5x** (specification effort)

# Outline

Push-button

Pay-as-you-go


Usage scenario

Full-stack


Evaluation [ Practical | No performance overhead ]

[vigor.epfl.ch](http://vigor.epfl.ch)

# Verification Effort



# Verification Effort



# Specification Example


```
1  from state import flow_emap, int_devices
2  EXP_TIME = 10 * 1000
3  EXT_DEVICE = 1
4
5  if a_packet_received:
6      flow_emap.expire_all(now - EXP_TIME)
7
8  h3 = pop_header(tcpudp, on_mismatch=[[[], []]])
9  h2 = pop_header(ipv4, on_mismatch=[[[], []]])
10 h1 = pop_header(ether, on_mismatch=[[[], []]])
11
12 if received_on_port == EXT_DEVICE:
13     internal_flow = FlowIdc(h3.dst_port, h3.src_port, h2.daddr, h2.saddr, h2.npid)
14     if flow_emap.has(internal_flow):
15         fl_id = flow_emap.get(internal_flow)
16         flow_emap.refresh_idx(fl_id, now)
17         out_port = vector_get(int_devices, fl_id)
18         return ([out_port], [ether(h1, saddr=..., daddr=...), ipv4(h2, cksum=...), tcpudp(h3)])
19     else:
20         return ([], [])
21 else:
22     internal_flow = FlowIdc(h3.src_port, h3.dst_port, h2.saddr, h2.daddr, h2.npid)
23     if flow_emap.has(internal_flow):
24         fl_id = flow_emap.get(internal_flow)
25         flow_emap.refresh_idx(fl_id, now)
26     else:
27         if not flow_emap.full():
28             fl_id = the_index_allocated
29             flow_emap.add(internal_flow, fl_id, now)
30             vector_set(int_devices, fl_id, received_on_port)
31     return ([EXT_DEVICE], [ether(h1, saddr=..., daddr=...), ipv4(h2, cksum=...), tcpudp(h3)])
```

# One-Off Property

```
1 if received_on_port == EXT_DEVICE and flow_emap.has(internal_flow):
2     fl_id = flow_emap.get(internal_flow)
3     flow_emap.refresh_idx(fl_id, now)
4     out_port = vector_get(int_devices, fl_id)
5     return ([out_port], [ether(h1, saddr=..., daddr=...), ipv4(h2, csum=...), tcputdp(h3)])
6 else:
7     pass
```


# Pay-as-you-go: ‘pass’

```
1 if received_on_port == EXT_DEVICE and flow_emap.has(internal_flow):
2     fl_id = flow_emap.get(internal_flow)
3     flow_emap.refresh_idx(fl_id, now)
4     out_port = vector_get(int_devices, fl_id)
5     return ([out_port], [ether(h1, saddr=..., daddr=...), ipv4(h2, cksm=...), tcputdp(h3)])
6 else:
7     pass
```



# Pay-as-you-go: ‘...’(ellipsis)

```
1 if received_on_port == EXT_DEVICE and flow_emap.has(internal_flow):
2     fl_id = flow_emap.get(internal_flow)
3     flow_emap.refresh_idx(fl_id, now)
4     out_port = vector_get(int_devices, fl_id)
5     return ([out_port], [ether(h1, saddr=..., daddr=...), ipv4(h2, cksum=...), tcpudp(h3)])
6 else:
7     pass
```



# Outline

Push-button

Pay-as-you-go

Usage scenario

Full-stack

Evaluation [ Practical | No performance overhead ]

[vigor.epfl.ch](http://vigor.epfl.ch)

# Usage Scenario

1. Implement, using Vigor library
2. Get the (standard) spec
3. ???
4. VERIFIED

# Usage Scenario

1. Implement, using Vigor library
  2. Get the (standard) spec
  3. 

```
vigor$ make verify
```
  4. VERIFIED
- Counterexample (execution trace)
- 

# Outline

Push-button

Pay-as-you-go

Usage scenario

Full-stack

Evaluation [ Practical | No performance overhead ]

[vigor.epfl.ch](http://vigor.epfl.ch)

# Recap: Insights and Design

Observations:

- NFs naturally split into NF logic + data structures
- NF specs and code use common abstractions
- Writing full specs is costly
- Most code in SW stack has no impact on desired properties

How Vigor leverages them:

- Verify different parts with different techniques, then stitch
- Adopt those abstractions for the Vigor library API and contracts
- Support incremental specs
- Exclude irrelevant code and verify the rest

# Outline

Push-button

Pay-as-you-go

Usage scenario

Full-stack

Evaluation [ Practical | No performance overhead ]

[vigor.epfl.ch](http://vigor.epfl.ch)

# Vigor NFs

- NAT
- Load balancer (Maglev)
- Traffic policer
- MAC-learning bridge
- Firewall

# Verify On Every Commit

NF	NF-only	Full-stack
NAT	<u>7 sec +</u> <u><math>54 \times 88</math> sec</u>	8 min + $488 \times 88$ sec
Load Balancer	23 sec + $146 \times 219$ sec	26 min + $1,336 \times 219$ sec
Traffic Policer	14 sec + $37 \times 82$ sec	6 min + $309 \times 82$ sec
Bridge	7 sec + $69 \times 80$ sec	10 min + $601 \times 80$ sec
Firewall	6 sec + $43 \times 88$ sec	7 min + $369 \times 88$ sec

- Verify NF in minutes
- Verify Full-stack in hours
- Embarrassingly parallel
- Found bugs in:
  - DPDK pkt proc framework
  - Intel NIC driver
  - Our NFs

# Concise and incremental specs

- **1 day** to formulate a full spec
- **30-60 LOC** size of a full spec
- **1-13 LOC** size of a minimal property
- **4-9 LOC** size of a property increment

# Outline

Push-button

Pay-as-you-go

Usage scenario

Full-stack

Evaluation [ Practical | No performance overhead ]

[vigor.epfl.ch](http://vigor.epfl.ch)

# Latency is Competitive


Latency

8  $\mu$ s


5  $\mu$ s

3  $\mu$ s

0  $\mu$ s



# Throughput is Competitive



# Recap: Eval

- Fast: can be done on every commit
- Easy: properties take just a few lines
- No performance overhead (modulo batching support)

Vigor enables NF developers to verify the  
Full stack of their NFs in a  
Pay-as-you-go manner with  
no verification expertise or effort  
(Push-button)



Vigor enables practical NF verification

Get it online! [vigor.epfl.ch](http://vigor.epfl.ch)

